Decision table — workload → best fit

If your workload is…	Reach for…
Large sequential training / checkpoints	Parallel FS (Lustre / EXAScaler)
Billions of small files	Object (S3-style)
Ultra-low latency / KV-cache serving	NVMe / NVMe-oF tier · object KV-native
Metadata-heavy workloads	Architecture matters more than raw bandwidth

If you remember one page, make it this one. Lab 2 (striping) makes the top row concrete.