Systems matrix — patterns, not products

System	Who / what	When	Why	How (in one line)
NFS	simplicity, ubiquity	home dirs, tools, shared config	everywhere; zero friction	single server exports a tree
GPFS / Spectrum Scale	balanced throughput + enterprise features	mixed enterprise HPC	features + performance	distributed block + wide metadata
Lustre / EXAScaler	maximum parallel throughput, striping	training data, multi-TB checkpoints, large sequential I/O	raw sequential bandwidth	MDT metadata + many OSTs, striped
Disaggregated (DDN Infinia et al.)	object-native, metadata-rich, multi-protocol	GPU-era access patterns	built for how AI reads data	scale-out nodes, every client → every node

EXAScaler = a hardened Lustre distribution. Each system is an answer to a different era's question — choose by workload, not brand.